

Rhea Madhogarhia

Essay for the Ann W. and Spencer T. Olin-Chancellor's Fellowship

For the past three years, my intellectual life has been defined by a tension between the "group mind" of improvisational comedy and the "black box" of machine learning. As a director, trainer, and improviser with the improvisational and sketch comedy group Off-Off Campus, I learned that a scene only works if the players maintain a shared, grounded reality. We call this "world-building." If I hand my scene partner an invisible cup, they must drink from it. The humor and humanity emerge from our mutual commitment to a truth we are discovering together.

Conversely, as a computer scientist, I have studied and used systems that generate "reality" without sharing it. Large Language Models (LLMs) are the ultimate improvisers, capable of generating fluent, human-like text. Yet, unlike my scene partners, they lack our "world." This contrast between the grounded connection of human improvisation and the stochastic mimicry of artificial intelligence is the driving force behind my decision to pursue graduate study in Philosophy-Neuroscience-Psychology (PNP).

In the summer of 2024, I worked as a Machine Learning Engineering Intern at Reddit. My task was to build a signal-driven First Pass Ranker to improve search relevance for millions of users. Technically, I succeeded; using Bayesian Optimization, I tuned the model to generate statistically significant improvements in quality, not just quantitative, user engagement. However, philosophically, I was unsettled. I realized I could alter global information flows by tweaking parameters that optimized for "engagement" without *fully* understanding the semantic implications of the content I was promoting. I had trust in our metrics, but that trust didn't feel generalizable to most projects.

This experience exposed an anxiety that we are integrating powerful AI into human society without the tools to audit its internal reasoning. I realized that Computer Science alone could not solve this problem. Philosophy helps define "understanding", and the tools of Neuroscience help map the biological mechanisms of thought. I wanted to leave the industry track to pursue the rigorous, interdisciplinary inquiry offered by the PNP PhD program.

However, the resilience to pursue this path, and shift away from a desire for a lucrative tech job that was just thrilling "enough" academically, came from a very different source. In Spring 2025, I took a leave of absence from the University of Chicago to manage a long-standing compulsive disorder. During this time, I stepped away from the high-pressure environment of academic tech and moved up a summer plan to participate in a work exchange in rural Central Java, Indonesia. Working as a conversational facilitator and English teacher at Miss Gati's Little IFE school in Temanggung, I found myself stripped of my usual digital tools and academic jargon. Teaching English to students with whom I shared no native language required a different kind of "mechanistic interpretability." I had to rely on shared attention, physical cues, and emotional resonance to build meaning from scratch. It was, in essence, the purest form of the "group mind" I had studied in improv. Connecting with children whom I couldn't speak to fluently was difficult, but it showed me something about connection and understanding that didn't need advanced linguistic fluency.

This lived experience in Indonesia reinforced my conviction that true intelligence is socially grounded and embodied—traits that current AI lacks. It also taught me resilience and humility. Stepping off the traditional "high-achiever" treadmill common to many UChicago students to prioritize my well-being was a difficult decision, but it gave me the clarity to return to academia not just as a student of computer science, but as a student of the human condition and the social sciences (something I am so grateful UChicago's core curriculum has provided me. I am returning to my studies with a renewed focus, eager to dissect the systems we are building to ensure they align with human values. For example, I'm currently taking Machine Learning, Philosophy of AI, Linguistic Truth, and Machine Learning for Computer Systems for my Fall 2025 coursework.

My work is inherently a bit liminal. To computer scientists, I am the researcher asking "philosophical" questions about intent; to philosophers, I am the engineer demanding "mechanistic" proof in neural weights. I faced this issue when looking for an advisor for my undergraduate cognitive science thesis, which is very interdisciplinary in nature. I am applying to the Olin-Chancellor's Fellowship because it is specifically designed for scholars who straddle these divides.

I view the OCF cohort as an interdisciplinary "group mind." In improv, the "Yes, And" principle suggests that the best ideas come from accepting and building upon the diverse offers of others. I want to bring this mindset to the OCF community. I hope that I can encourage the OCF community to engage in informal, improv-style gatherings or open discussions where fellows from disparate fields (a historian, a biologist, etc) tackle a single modern problem, such as "Memory" or "Truth." My goal is to participate in a community that actively butts heads in search of collaboration.

Furthermore, I want to demystify the "black box" of AI for the general public, helping to translate high-level technical risks into accessible language for policymakers and community leaders. My involvement in non-profit organizations like The Rithm Project and the Young Futures Foundation has laid the groundwork for my interest in tech policy and digital literacy, and I hope my research empirically grounds my opinions and interest in these separate fields.

We are entering a time where artificial agents act as tutors, therapists, and judges. I hope to remain dedicated to developing the "causal auditing" frameworks necessary to determine when these models are safe and appropriate to use. This could be working as a post-doctoral researcher in a lab, working in safety and alignment positions at institutions like Google DeepMind or Anthropic, teaching, or starting my own lab. There are many pathways I'd like to keep open for post-grad, but I want to ensure that as we outsource more of our cognitive labor to machines, we do so with a scientific understanding of exactly what they are and what they are not.

Word Count: 999