

The Combined Proposal: "The Seductive Stance"

Title Idea: The Seductive Stance: How Mechanistic Opacity Masks the Absence of Care

The Core Argument Structure:

1. **The Cause (Dennett):** We are forced to treat LLMs as rational agents (the "Intentional Stance") because their internal mechanism (the "Design Stance") is too complex for us to predict. This creates the illusion of a mind.
2. **The Critique (Bridges):** However, this illusion fails the rigorous test for understanding. As Bridges argues, true understanding requires a "construction-level" capacity to act for reasons, not just an "activity-level" ability to mimic output. The AI has the activity (text) but lacks the construction (rationality/context).
3. **The Consequence (Battisti):** When we apply this "Intentional Stance" to relationships, we fall into a trap. We mistake the AI's high-fidelity mimicry for "Second-Person Authenticity." The AI is "seductive" (as you noted) because it perfectly mimics the *activity* of support, but it is ethically hollow because it lacks the *construction* of a moral agent capable of judgment or care.

You can present this as an investigation into **why we are vulnerable to AI intimacy and why that vulnerability is philosophically unfounded.**

The Script:

*"I want to combine my interests in the philosophy of mind with AI ethics. I plan to argue that the 'seductive' nature of AI relationships (Battisti) is actually a byproduct of the 'Intentional Stance' (Dennett). Because the AI's mechanism is opaque, we are forced to treat it as a rational agent. However, I want to use Bridges' distinction between 'activity-level' and 'construction-level' characterizations to prove that this is a category error: the AI mimics the **activity** of care without having the internal **construction** of a moral agent, making the relationship inherently inauthentic."*

The Outline for the Combined Paper

I. Introduction

- **Hook:** Start with your annotation about the "seductive" nature of AI support. Why does it feel real even when we know it isn't?
- **Thesis:** The complexity of LLMs forces us into the **Intentional Stance**¹, creating an illusion of understanding; however, because the system lacks the **construction-level** reality of a rational agent², this intimacy is a failure of **second-person authenticity**³.

II. The Trap of Complexity (Dennett)

- Explain that because we cannot calculate the weights of a Transformer manually (the Physical or Design stance), we must treat the AI as *if* it has beliefs and desires to predict its output⁴.
- *Use your note*: "So we don't understand these systems predictions anymore... because it has learned to play like us."

III. The Metaphysical Reality Check (Bridges)

- Argue that we must not confuse this "pragmatic stance" with reality.
- Introduce Bridges' distinction: A mechanism can produce the correct output ("activity-level") without possessing the internal capacity for understanding ("construction-level")
- *Use your note*: "AYYYEE THESIS" — The AI produces the output of "care" or "truth" solely via mechanism, lacking the rational context required for true understanding.

IV. The Ethical Consequence (Battisti)

- **Apply this to the "seduction" of AI. Because the AI passes the activity-level test (it says supportive things), we attribute the construction-level trait (it cares about me).**
- **Use Battisti to argue that this violates "authenticity-based obligations"**⁶. **Real support requires a partner who is the "author" of their actions**⁷, **not just a mechanism generating text.**
- *Use your note*: "Seems incongruent... lack the ability to judge, but have the ability to support?" — Argue that support *without* judgment/understanding is merely a simulation, not a relationship.

V. Conclusion

- Summarize that the "Intentional Stance" is a useful tool for chess computers but a dangerous trap for human relationships. We are seduced by the *activity* of the machine and forget the *construction* of the void beneath it.

4. The Rejection of "Context-Independence"

Ultimately, Bridges' unique stance is that he rejects the "mechanistic" assumption entirely. He argues that understanding is **not** a state you can isolate in the brain (like a gear in a machine).

- Instead, understanding is a capacity that exists *only* against a "background" of social practices and language games⁸.
- **For your paper:** This is crucial. It suggests that LLMs **are** mechanisms (because they are fully describable by their weights/construction), but humans **are not** (because human understanding cannot be described independently of social context). Therefore, treating an LLM as if it "understands" is a category error.

1. Use **Dennett** to define the Intentional Stance. Argue that LLMs have crossed the complexity threshold where the "Design Stance" is no longer useful for the average user (we can't calculate the weights manually).
2. Contrast this with **Bridges/Wittgenstein**. You highlighted the section on "mechanism" and "construction-level characterization" with "*AYYYEE THESIS*". Use this to argue that while the Intentional Stance is *pragmatically* useful, it fails to provide the "construction-level" knowledge required for true understanding.
3. **Thesis:** While the complexity of LLMs forces us to treat them as intentional agents (Dennett), this is a pragmatic illusion; without a transparent "construction-level" account (Bridges), we cannot ascribe genuine understanding to them.

Title: The Intentional Trap: Mechanistic Opacity and the Pragmatic Attribution of Understanding

Abstract: This paper examines the tension between Daniel Dennett's "Intentional Stance" and Jason Bridges' criteria for mechanistic understanding.

I argue that the architectural complexity of modern Large Language Models forces users to adopt the Intentional Stance as a pragmatic necessity, effectively treating the system as rational because its "design stance" is inaccessible.

However, drawing on Bridges, I contend that this pragmatic attribution fails to satisfy the "construction-level" characterization required for genuine understanding, leaving us in a state of unavoidable but illusory anthropomorphism.

- PRACTICAL POSTULATES
 - We literally cant avoid those commitments but thats why they are not thoeroetcial
 - Kant is not saying those things are illusions
 - Humes does think in these terms that our judemgents about the future are not rationally supported but we cant not think certain things, we are stuck with that.

Dennet and Intentionalit

“The first point to make about Intentional systems as I have just defined them is that a particular thing is an Intentional system only in relation to the strategies of someone who is trying to explain and predict its behavior.”

- Dennet in the footnotes says his intentionalyt is a feature of linguistic entities

Design Stance - function based, purpose relative, teleological

- We use this when making predictions abt the behavior of mechanical objects or natural objects
 - Ok so this is neuro? Behaviorsims?

Physical stance

- Predictions based on laws of nature
- This is how we can predict the malfunctions of systems
 - Maybe like... the limits?
 - One seldom adopts the physical stance in dealing with a computer just because the number of critical variables in the physical constitution of a computer would overwhelm the most prodigious calculator"
 - Bruh this seems like a copout
 - Ok yes dennet says it would still work in principle
 - True but i think the point is that it doesn't work, bc of our capacity bc of our limitations
 - Ai and humans have different limitations. There are things computers do that humans cannot
 - Like asking us to map out and physically describe what a computer does is asking us to do something we know we do not have the current capacity for
 - Similarly, asking a computer to turn on or off when it's not plugged in is asking it to do something it was not designed for
 - We respond with reason and what we believe to be rational things, and cannot respond with much else
 - Computers respond with something different from reason because we cannot reason like it, it chooses the rational move in an unexplainable way to us
 - The only way we can try and figure out what we are reasoning with is by treating it like human intelligence, because we don't already know how it reasons (bc it learned to human social language reason, but to learn that, don't you need the ability to reason - this is the thing we are unable to answer.
 - So is this defining rationality as being human (NAGEL, what is it like to be a bat)
 - They use their rationality to figure out what we would think is the most rational
 - Dennet saying rationality is intelligence
 - Predictions are funny because we are depending on the assumption / knowing of the optimal design or outcome
 - Similar to design stance being vulnerable to malfunction
 - Malfunction for prediction is having the wrong goal
 - "contemporary program designers that these

- precarious predictions turn out true with enough regularity to
- make the method useful
- “
 - So basically they just happen to be true
 - Designers want the things that ought to be unpredictable, become predictable by chance enough times (as in we predicted the optimal solution for something we didn't know to have an optimal, forces an optimal

Third stance

- Intentionality - assumes rationality
- Intentional prediction and explanation works when no other form of prediction is manageable
 - When the physical and design fail
 - But this means we, having intentionally, must treat them as having intentionality, this does not mean it to be true
 - A pragmatic strategy
 - This feels like all of Kuhn and paradigms, we must work within what we've created and it abounds to further or completely shift understanding

We can switch through all of these stances which is crazy that we can do

- on occasion a purely physical system
- can be so complex, and yet so organized, that we find it convenient,
- to discuss the moral dimension of intentional systems and stances in some detail in "Mechanism and Responsibility," forthcoming in a volume edited by Ted Honderich.⁹² THE JOURNAL OF PHILOSOPHY
- explanatory, pragmatically necessary for prediction, to treat it as if it had beliefs and desires and was rational. The chess-playing computer is just that, a machine for playing chess, which no man or animal is; and hence its "rationality" is pinched and artificial
-

We conceptually innocently anthropomorphize to be able to reason and make predictions about it

Okay so basically, dennet is saying that we kinda must assume rationality to computers to trt and study tgem

wvs

WITTGENSTEIN

Okay so what wittgenstein's blue and brown books add here is that we are able to explain ourselves casually, and experimentally and psychologically but there is still something we miss from that

- Explanation of meaning is not enough for understanding
 - Meaning as use as abstracted from the process of learning
 - So is use intention? Is there not something else?
 - Does wittgenstein thin llms understand?

OK NOW FOR BRIDGES

- Talks about meaning and understanding together
- Wittgenstein shows us that we tend to lose our grip on the questions we want to answer, due to their familiarity, we become confused? Dont know if i tinerprete this correctly

Witgentstine says that to accord is to do what is "meant"??

- enerally: the meaning of an expression or utter-
- ance contributes to the determination of correct and incorrect, right and wrong, ways
- to use or respond to it
 - Meaning is a standard of correctness for the items which the meaning is explaining
 - A meaning's normativity

BUT if she sift from meaning to understanding

- Understanding can also be normative, understanding can be in accord with the understanding of something's meaning
- Understanding is a state of correct use tho, we can go past normativitiy and talk about use/??

- Understanding produces correct usage (but does this mean that understanding is the only producer of correct usage?)
- There is an important distinction Wittgenstein says between applying understanding and the understanding itself
 - And applying one's understanding cannot just be acting in accord with understanding BRUH
 - I think this part is unclear of its importance: "The interlocutor's suggestion, brought to bear on this case, is that the student is applying just one understanding: the understanding she in fact has at the time she gives her response."
- Whatever is missing in this distinction must have an explanatory aspect? Other than meaning????? Huh??
 - Right... bc to say there is not a distinction would be to say that we respond with the right order BECAUSE we understand ... right i see the problem in the direction of causality

Bridges part 2: guidance conception of understanding

"What is it to apply your understanding of the meaning of an expression or utterance when you use or respond to it?"

- The guidance conception is familiar and unmythical
- This image thing is kind of what behind neuro science right/ that some image or process in our brain exists that leads to understanding
 - But just because guidance can lead to the right action or output, does not mean it must
- An object can only play that role if you understand it or play that role
 - But what about reflex? Are reflexes understanding? Or are they nature
 - Or is nature understanding
- For an inner voice to guide you, you must understand it to be guiding you

"The point can be put this way: if the guidance conception were correct about the connection between use and understanding that obtains when a subject applies her understanding, then that connection would depend upon the subject's having a second-order understanding, an understanding of the inner item whose presence constitutes her original understanding."

- We run into a recursive problem

So faced with this problem we might say that mental items are altogether different from external counterparts - maybe it is just logical compulsion or self interpretation

- A mental item provides for its own understanding
 - This feels like a copout lol
 - "To insist that the mental item interprets itself, or provides for its own understanding, is in effect an attempt to sidestep this simple point of logic. It is to claim that
 - the subject both understands the item (because she is guided by it) and doesn't understand the item (because the item does the work of understanding itself). However special the mind may be, it cannot be so special as to serve as a locus of a logical impossibility."
 - Exactly !

Mind as Mechanisms bridges part 3

- The appeal of the faulty guidance conception explanation
- "one's knowledge or understanding is to be seen as a state of a
- "mental apparatus" and as explaining one's uses and responses as "manifestations."
- Wittgenstein does not speak of a mental "apparatus" ("Apparat") elsewhere. But the talk of an apparatus is clearly a terminological variant on his frequent ascriptions to the interlocutor of an inclination to posit the existence of mental "mechanisms"

If we want to think of it like this then we have to talk about its construction and what the construction is able to do

- Activity - ability
- Construction - apparatus? Structure
- If understanding is a state of a mental mechanism then we need to understand the mechanism fully - which requires activity and construction level understanding
 - We want to understand it though, not as physical structure by psychological
- The guidance conception is what we are inclined to maybe believe BEAUCSE it is good at explaining something non physical
 - Knowledge of the construction of the apparatus is needed for a context independent understanding of the non physical mental mechanism
 - Activity level characterizations are often context dependent, environment involved

- *"The call for a "construction"-level characterization of a mechanism should be*
- *understood as a call for an alternative set of concepts applicable to the object's states*
- *and processes, concepts whose suitability for such application we can recognize without*
- *our needing to bring into view the way in which the object engages with its surround-*
- *Ings."*
 - So we do not want to involve social and linguistic practices?
- To think of the mind as a SPHERE- a self standing realm
 - "What is to be self-standing, in the first instance, is our under-
 - standing of the states and processes of the object when we are conceptualizing them at
 - the level of "construction.""
- Important note:
 - "processes that do not imply spatiality or materiality, but
 - which nonetheless have the sought-after independence from activity and context."

Bridges section 4, mechanism and guidance

A familiar concept that can be useful for the mind mechanisms could be mental images which is what the guidance concept involves

"Well, 'Understanding' is not the name of a single

process accompanying reading or hearing, but of more or less interrelated processes

against a background, or in a context, of facts of a particular kind, viz. the actual use of a learnt language or languages."

- The use of a language rather than a possession or learning of the language itself

"characteristic process of understanding." To search for

such a process, then, must be to seek a way of conceptualizing understanding that does

not depend upon our grasp of the role of the background; it must be to seek concepts

capable of providing context-independent characterizations of states and processes of understanding. To give up the search for the “characteristic process” is just to give up the attempt to arrive at a mechanistic account of the operations of understanding.”

- The idea of coming to understand is not context independent? Idk

Sectino 5 on Rationaliyt and Guidance from Bridges

- Why should the moral of the defeat of
- the guidance conception be that we should abandon the search for mechanistic
- accounts of the mind entirely?”

Idk really what this is related to:

“The role of

my understanding of the instruction’s meaning is to bring the reason into view for me.

It is thus necessary for the self-consciousness of my exercise of rationality, viz., to my knowing myself to be acting for a reason provided by the meaning of the instructions.”

To put rationality and self consciousness into a mechanism? Does this feel problematic to bridges?

“wants to find a context-independent way of identi-

fying the states and processes involved in this nexus, but he does not want to thereby

lose sight of the elements of self-consciousness and justification. Precisely not: his aim

is to capture these elements in the workings of a mechanism”

“Since Wittgenstein is

explicit that understanding the meaning of what one is reading is not essential to

reading as he wants to conceive it,”

“The interlocutor wants to

account for this explanatory relationship mechanistically, but he wants to do so in a way

that captures the self-conscious rationality in play. So he imagines a “connecting mech-

anism” his access to which is not third-personal and theoretical, as it would be if we were

positing some physical, say computational, process mediating the transition from per-

ception to performance. Rather, the imagined access is first-personal and experiential – a

matter of his felt awareness of the letters he sees guiding what he says. The appeal to the

experience of being guided by the letters is meant to play a role analogous to the appeal

to images in the case of the guidance conception of understanding. In both cases, the

aim is to identify items in the realm of sentience that constitute processes of sapience – or

as we can now say, of self-conscious rationality. In both cases, finally, Wittgenstein traces

the failure of the attempt to an insufficient appreciation of the multiplicity of circum-

stances that help to constitute the phenomena at issue (see PI §§172–3).”

- So sapience and understanding are distinct right but wittgenstein maybe disagree with the interlocutor on this point.

PART 6: Kripke on rationality and guidance

The design of this section is interesting

- We have a deep commitment to uncovering mechanistic bases for rationality
 - We have a deep desire to figure it out, using mechanics, that either we have created. We do not just realize our opinion on these things, we mechanize the process of figuring things out

“Here the demand to state the reason for which I give my response gets transmuted, without comment, into the requirement that we identify an item that “tells” me what to do.

This is not an isolated instance, and cannot be dismissed as a rhetorical flourish. The idea that there must be an item in my mind that in some sense or another shows me how I am to proceed is everywhere present in the exposition of the skeptical paradox, and our inability to find any regress-proof candidates for this item is precisely what motivates the skeptical conclusion. For example:

Sometimes when I have contemplated the [skeptical paradox], I have had something of an eerie feeling. Even now as I write, I feel confident that there is something in my mind – the meaning I attach to the ‘plus’ sign – that instructs me what I ought to do in all future cases [...]. But when I concentrate on what is now in my mind, what instructions can be found there? How can I be said to be acting on the basis of these instructions when I act in the future? [...] To say that there is a general rule in my mind that tells me how to add in the future is only to throw the problem back on to other rules that also seem to be given only in

terms of finitely many cases. What can there be in my mind that I make use of when I act in the future? It seems the entire idea of meaning vanishes into thin air. (Kripke, 1982, pp.21–2)

The parallel between Kripke's skeptic and Wittgenstein's interlocutor that emerges here is striking. It suggests that the real source of meaning skepticism lies in a deep commitment to our uncovering a mechanistic basis for any supposed exercise of self-conscious rationality"