

Reading Response 3: Rideshare Transparency: Translating Gig Worker Insights on
AI Platform Design to Policy – Nagaraj et. al

Summary: (made sure to separate into sections as said in class)

Goals: This study investigates how the lack of transparency in AI-driven rideshare platforms creates information asymmetries that harm gig workers. They mentioned that prior studies rely on speculative design, while this paper evaluates existing platform designs to identify transparency gaps. They also fill the gap in existing literature for actionable steps by proposing policy-driven "Rideshare Transparency Reports".

Takeaways: The authors combined interviews with LLM-based analysis of Reddit comments. They identified four areas where opacity of AI-driven rideshare interfaces causes financial, physical, and psychological harm. The four are promotions/misleading gamification, fares/unpredictable earnings, routes/safety concerns, and task allocation/fears of bias/discrimination. The study finds that drivers attempt to avoid the consequences of rideshare app algorithms through strategies like "multiapping" or auto-accepting rides. **Strengths:** A major strength is the multi-method approach, which combines qualitative interviews with the quantitative breadth of massive amounts of community data from Reddit. This allows the researchers to validate their qualitative interview themes against popular issues shared across internet communities of drivers. Also, the paper offers policy frameworks - the Transparency Report. **Weaknesses:** The reliance on Reddit data introduces biases (male, tech-savvy, English-speaking users) that do not represent the diversity of the gig economy. Also, while they validated the LLM's output, they were using a general-purpose LLM and were asking the LLM to categorize the data into predefined categories. Potentially (although not necessarily), this forces some data to fit into categories that the data does not belong in. **Limitations:** The researchers included a positionality statement that acknowledged that their research is likely impacted by researcher bias. The interview sample size was only 9, and geographically limited to the Northeast US. The study is also US-centric and limited to English-language forums, excluding the perspectives of non-English speaking drivers who are probably the most vulnerable to algorithmic exploitation.

Reaction:

Interesting Parts: As a student interested in the "black box" problem (discussing this in my cognitive science thesis), I found the methodology and justification for methodology a bit ironic. The researchers used a general-purpose LLM to study the harms of opaque AI features in rideshare platforms. To me, this seems a bit silly... Can we trust a black box to audit or qualify another black box? LLMs do not have a commitment to the truth, which has made me quite skeptical of using LLM annotations officially in research. **Applications:** The "Rideshare Transparency Report" is a compelling application of AI governance, a field I've become increasingly interested in as a computer science student and industry engineering intern. It shifts the focus from "explainability" or putting the terribly difficult burden of making these systems more translucent, to "accountability" via data disclosure. Just knowing how an LLM works and where it is being used can be enough to deter people from using AI in scenarios they wouldn't normally be comfortable using it in.